

CSCI 2021: x86-64 Control Flow

Chris Kauffman

*Last Updated:
Mon Mar 4 11:41:59 CST 2019*

Logistics

Reading Bryant/O'Hallaron

- ▶ Ch 3.6: Control Flow
- ▶ Ch 3.7: Procedure calls

Goals

- ▶ Jumps and Control flow
x86-64
- ▶ Procedure calls

Assignment 3: Post Today

- ▶ Problem 1: Thermometer
Assembly Functions (50%)
- ▶ Problem 2: Binary Bomb via
debugger (50%)

Date	Event
Mon 3/4	Basics Wrap-up
Tue/Wed	Assembly Lab
Wed 3/6	Control Flow
Fri 3/8	Procedures
Mon 3/11	Assembly Wrap
Tue/Wed	Review Lab
Wed 3/13	Review A3 Due
Fri 3/14	Exam 2

Control Flow in Assembly and the Instruction Pointer

- ▶ No high-level conditional or looping constructs in assembly
- ▶ Only `%rip`: Instruction Pointer or "Program Counter": memory address of the next instruction to execute
- ▶ Don't mess with `%rip` by hand: automatically increases as instructions execute so the next valid instruction is referenced
- ▶ Jump instructions modify `%rip` to go elsewhere
- ▶ Typically label assembly code with positions of instructions that will be the target of jumps
- ▶ **Unconditional Jump** Instructions always jump to a new location.
- ▶ **Comparison / Test** Instruction, sets EFLAGS bits indicating relation between registers/values
- ▶ **Conditional Jump** Instruction, jumps to a new location if certain bits of EFLAGS are set, ignored if bits not set

Examine: Loop Sum with Instruction Pointer (rip)

- ▶ Can see direct effects on rip in disassembled code
- ▶ rip increases corresponding to instruction length
- ▶ Jumps include address for next rip

```
// C Code equivalent
int sum=0, i=1, lim=100;
while(i<=lim){
    sum += i;
    i++;
}
return sum;
```

000000000000005fa <main>:

ADDR	HEX-OPCODES	ASSEMBLY	EFFECT ON RIP
5fa:	48 c7 c0 00 00 00 00	mov \$0x0,%rax	# rip = 5fa -> 601
601:	48 c7 c1 01 00 00 00	mov \$0x1,%rcx	# rip = 601 -> 608
608:	48 c7 c2 64 00 00 00	mov \$0x64,%rdx	# rip = 608 -> 60f

0000000000000060f <LOOP>:

60f:	48 39 d1	cmp %rdx,%rcx	# rip = 60f -> 612
612:	7f 08	jg 61c <END>	# rip = 612 -> 614 OR 61c
614:	48 01 c8	add %rcx,%rax	# rip = 614 -> 617
617:	48 ff c1	inc %rcx	# rip = 617 -> 61a
61a:	eb f3	jmp 60f <LOOP>	# rip = 61a -> 60f

0000000000000061c <END>:

61c:	c3	retq	# rip 61c -> return address
------	----	------	-----------------------------

FLAGS: Condition Codes Register

- ▶ Most CPUs have a special register with "flags" for various conditions
- ▶ In x86-64 this register goes by the following names

Name	Width	Notes
FLAGS	16-bit	Most important bits in first 16
EFLAGS	32-bit	Name shown in gdb
RFLAGS	64-bit	Not used normally

- ▶ Bits in FLAGS register are **automatically** set based on results of other operations
- ▶ Pertinent examples with conditional execution

Bit	Abbrev	Name	Description
0	CF	Carry flag	Set if last op caused unsigned overflow
6	ZF	Zero flag	Set if last op yielded a 0 result
7	SF	Sign flag	Set if last op yielded a negative
8	TF	Trap flag	Used by gdb to stop after one ASM instruction
9	IF	Interrupt flag	1: handle hardware interrupts, 0: ignore them
11	OF	Overflow flag	Set if last op caused signed overflow/underflow

Comparisons and Tests

- ▶ Set the EFLAGS register by using comparison instructions

Name	Instruction	Examples	Notes
Compare	<code>cmpX B, A</code>	<code>cmpl \$1,%eax</code>	Like <code>if(eax > 1){...}</code>
	Like: <code>A - B</code>	<code>cmpq %rsi,%rdi</code>	Like <code>if(rdi > rsi){...}</code>
Test	<code>testX B, A</code>	<code>testq %rcx,%rdx</code>	Like <code>if(rdx & rcx){...}</code>
	Like: <code>A & B</code>	<code>testl %rax,%rax</code>	Like <code>if(rax){...}</code>

- ▶ B,A are NOT altered with `cmp` and `test` instructions
- ▶ Only the stat
- ▶ Immediates like `$2` must be the first argument B
- ▶ `cmpX` and `testX` will alter EFLAGS register so that certain bits are set indicating equality, less than, greater than, etc.

```
cmpl $1, %eax # compare: eax > 1 VIA eax - 1
## EFLAGS bits set based on result of eax - 1
##   ZF (zero flag) now 1 if eax==1
##   SF (sign flag) now 1 if eax<1
```

```
testq %rax,%rax # test rax VIA rax & rax
## EFLAGS bits set based on result of rax & rax
##   ZF (zero flag) now 1 if rax==0 (falsey)
##   ZF (zero flag) now 0 if rax!=0 (truthy)
```

Jump Instruction Summary

<i>Instruction</i>	<i>Effect</i>
jmp LAB	Unconditional jump
je LAB	Equal / zero
jz LAB	
jne LAB	Not equal / non-zero
jnz LAB	
js LAB	Negative
jns LAB	Nonnegative
jg LAB	Greater-than signed
jge LAB	Greater-than-equal signed
jl LAB	Less-than signed
jle LAB	Less-than-equal signed
ja LAB	Above unsigned
jae LAB	Above-equal unsigned
jb LAB	Below unsigned
jbe LAB	Below-equal unsigned
jmp *OPER	Unconditional jump to variable address

- ▶ Synonyms exist for many of these
- ▶ ja used by compiler for `if(a < 0 || a > lim)`
Consider sign/unsigned to explain why
- ▶ `jmp %rdx` allows **function pointers**, powerful but no time to discuss

Examine: Compiler Comparison Inversion

- ▶ Often compiler inverts comparisons
- ▶ `i < n` becomes `cmpX / jge` (jump greater/equal)
- ▶ `i == 0` becomes `cmpX / jne` (jump not equal)
- ▶ This allows "true" case to fall through immediately
- ▶ Depending on structure, may have additional jumps
 - ▶ `if(){ .. }` usually has a single jump
 - ▶ `if(){ } else { }` may have a couple

```
## Assembly translation of
## if(rbx >= 2){
##   rdx = 10;
## }
## else{
##   rdx = 5;
## }
## return rdx;
    cmpq    $2,%rbx      # compare: rbx-0
    jl     .LESSTHAN    # goto less than
    ## if(rbx >= 2){
    movq    $10,%rdx     # greater/equal
    ## }
    jmp     .AFTER
.LESSTHAN:
    ## else{
    movq    $5,%rdx      # less than
    ## }
.AFTER:
    ## rdx is 10 if rbx >= 2
    ## rdx is 5 otherwise
    movq    %rdx,%rax
    ret
```


Exercise: Other Kinds of Conditions

Other Things to Look For

- ▶ `testl %eax,%eax` often used to check
- ▶ Zero/nonzero, followed by `je` / `jz` / `jne` / `jnz`
- ▶ Also works for NULL checks
- ▶ Negative Values, followed by `js` / `jns` (jump sign / jump no sign)

See `jmp_tests_asm.s`

- ▶ Trace the execution of this code
- ▶ Determine return value in `%eax`

cmov Family: Conditional Moves

- ▶ A family of instructions allows conditional movement of data into registers
- ▶ Can limit jumping in simple assignments

```
cmpq    %r8,%r9
cmovge   %r11,%r10  # if(r9 >= r8) { r10 = r11 }
cmovg    %r13,%r12  # if(r9 >  r8) { r12 = r13 }
```

- ▶ Note that condition flags are set on arithmetic operations
- ▶ cmpX is like subQ: both set FLAG bits the same
- ▶ Greater than is based on the SIGN flag indicating subtraction would be negative allowing the following:

```
subq     %r8,%r9      # r9 = r9 - r8
cmovge    %r11,%r10   # if(r9 >= 0) { r10 = r11 }
cmovg     %r13,%r12   # if(r9 >  0) { r12 = r13 }
```

Procedure Calls

Have seen basics so far:

```
call  PROCNAME  # call a function
        ## Pushes return address %rip onto stack adjusting

movl   $0,%eax   # set up return value
ret                    # return from function
        ## Pops old %rip off of stack adjusting %rsp
```

Need several additional notions

- ▶ Control Transfer?
- ▶ Where are arguments to functions?
- ▶ Return value?
- ▶ Anything special about the registers?
- ▶ How is the stack used?

Control Transfer with Basics

The `call` instruction does two things

1. Push the "caller" return address onto the stack
 - ▶ Calculate next `rip` for current function: return address
 - ▶ Grow stack: push return address so $\text{rsp} = \text{rsp} - 8$
2. Change the `rip` to be the first instruction of the "callee" function: execution starts in new function

The `ret` function undoes these

1. Pop the return address off the stack to become the new `rip`
2. Shrinks stack for popped address: $\text{rsp} = \text{rsp} + 8$




Figure: Bryant/O'Hallaron Fig 3.26 demonstrates `call`/return in assembly

Example: Control Transfer with call

BEFORE CALL

main: ...

```
0x555555554687 <+11>:      mov     $0x5,%esi
=> 0x55555555468c <+16>:      callq   0x55555555466a <sum_range>
0x555555554691 <+21>:      mov     %eax,%ebx
```

rip = 0x55555555468c -> call -> 0x555555554691

rsp = 0x7fffffff460

(gdb) stepi

AFTER CALL

sum_range:

```
=> 0x55555555466a <+0>: mov     $0x0,%eax
0x55555555466f <+5>: jmp     0x555555554676 <.TOP>
```

rip = 0x55555555466a

rsp = 0x7fffffff458 # pushed return address: rsp -= 8

(gdb) x/xg \$rsp

0x7fffffff458: 0x555555554691 # return address in main

Control Transfer with ret

BEFORE RET:

sum_range:...

0x555555554678 <+2>: jle 0x555555554671 <.BODY>

=> 0x55555555467a <+4>: repz retq

rip = 0x55555555467a -> return

rsp = 0x7fffffff458

(gdb) x/xg \$rsp

0x7fffffff458: 0x555555554691 # return address in main

(gdb) stepi

AFTER RET

0x555555554687 <+11>: mov \$0x5,%esi

0x55555555468c <+16>: callq 0x55555555466a <sum_range>

=> 0x555555554691 <+21>: mov %eax,%ebx

rip = 0x555555554691

rsp = 0x7fffffff460 # popped return address: rsp += 8

Stack Alignment

- ▶ According to the strict x86-64 ABI, must align `rsp` (stack pointer) to 16-byte boundaries when calling functions
- ▶ Will often see arbitrary pushes or subtractions to align
 - ▶ Always enter a function with old `rip` on the stack
 - ▶ Means that it is aligned to 8-byte boundary
- ▶ `rsp` changes must be undone prior to return

```
main:                                # enter with at 8-byte boundary
    subq    $8, %rsp                # align stack for func calls
    ...
    call    sum_range               # call function
    ...
    addq    $8, %rsp                # remove rsp change
    ret
```
- ▶ Failing to align the stack may work but may break
- ▶ Failing to "undo" stack pointer changes will likely result in return to the wrong spot : major problems

x86-64 Register/Procedure Convention

- ▶ Used by Linux/Mac/BSD/General Unix
- ▶ Params and return in registers if possible

Parameters and Return

- ▶ First 6 arguments are put into
 1. rdi / edi / di (arg 1)
 2. rsi / esi / si (arg 2)
 3. rdx / edx / dx (arg 3)
 4. rcx / ecx / cx (arg 4)
 5. r8 / r8d / r8w (arg 5)
 6. r9 / r9d / r9w (arg 6)
- ▶ Additional arguments are pushed onto the stack
- ▶ Return Value in rax / eax / ...

Caller/Callee Save

Caller save registers: alter freely

```
rax rcx rdx rdi rsi  
r8  r9  r10 r11
```

Callee save registers: must restore these on return

```
rbx rbp r12 r13 r14  
r15
```

Careful messing with stack pointer

```
rsp # stack pointer
```


Pushing and Popping the Stack

- ▶ If local variables are needed on the stack, can use `push` / `pop` for these
- ▶ `pushX %reg`: grow `rsp` (lower value), move value to top of main memory stack,
 - ▶ `pushq %rax`: grows `rsp` by 8, puts contents of `rax` at top
 - ▶ `pushl $25`: grows `rsp` by 4, puts constant 5 at top of stack
- ▶ `popX %reg`: move value from top of main memory stack to `reg`, shrink `rsp` (higher value)
 - ▶ `popl %eax`: move `(%rsp)` to `eax`, shrink `rsp` by 4

```
main:
    pushq    %rbp                # save register, aligns stack
                                   # like subq $8,%rsp; movq %rbp, (%rsp)
    call     sum_range           # call function
    movl     %eax, %ebp          # save answer
    ...
    call     sum_range           # call function, ebp not affected
    ...
    popq     %rbp                # restore rbp, shrinks stack
                                   # like movq (%rsp), %rbp; addq $8, %rsp
    ret
```

Local Arguments May be on the Stack

- ▶ Variables that are large or need addresses usually on the stack
- ▶ Compiler calculates location as $\text{rsp} + \text{offsets}$

```
// C Code: locals.c
int set_buf(char *b, int *s);
int main(){
    // locals re-ordered on
    // stack by compiler
    int size = -1;
    char buf[16] = "push it";
    ...
    int x = set_buf(buf, &size);
    ...
    return 0;
}
```

-----+-----+-----		
REG	VALUE	Name
-----+-----+-----		
rsp	#1024	top of stack
		during main
-----+-----+-----		
MEM		
...
#1031	h	buf[3]
#1030	s	buf[2]
#1029	u	buf[1]
#1028	p	buf[0]
#1024	-1	size
-----+-----+-----		

EQUIVALENT ASSEMBLY

main:

subq	\$24, %rsp	# grow stack, min size 20	8+24 = 32
leaq	4(%rsp), %rdi	# address of buf arg1	16-byte aligned
leaq	0(%rsp), %rsi	# address of size arg2	for func call
call	set_buf	# call function	8 from prev rip
movl	%eax,%r8	# get return value	24 from subq
...			
addq	\$24, %rsp	# shrink stack size	
movl	\$0,%eax	# return 0	
ret			

Historical Aside: Base Pointer rbp was Important

```
int bar(int, int, int);
int foo(void) {
    int x = callee(1, 2, 3);
    return x+5;
}
```

- ▶ 32-bit x86 / IA32 assembly used rbp as bottom of stack frame, rsp as top.
- ▶ Push all arguments onto the stack when calling changing both rsp and rbp
- ▶ x86-64 default: **NOT USED** making rbp a general purpose register

Old x86 / IA32 calling sequence: set both %esp and %ebp for function call foo:

```
    pushl %ebp                # modifying ebp, save it
    ## Set up for function call to bar()
    movl  %esp,%ebp          # new frame for next function
    pushl 3                   # push all arguments to
    pushl 2                   # function onto stack
    pushl 1                   # no regs used
    call bar                  # call function, return val in %eax
    ## Tear down for function call bar()
    movl  %ebp,%esp          # restore stack top: args popped
    ## Continue with function foo()
    addl 5,%eax               # add onto answer
    popl  %ebp               # restore previous base pointer
    ret
```