

CSCI 2021: x86-64 Assembly Extras and Wrap

Chris Kauffman

*Last Updated:
Mon Mar 11 12:04:39 CDT 2019*

Logistics

Reading Bryant/O'Hallaron

Skim the following on
Assembly/C

- ▶ Ch 3.8-3.9: Arrays, Structs
- ▶ Ch 3.10: Pointers/Security
- ▶ Ch 3.11: Floating Point

Goals

- ▶ Finish control
- ▶ Data in Assembly
- ▶ Security Risks
- ▶ Floating Point Instr/Regs

Date	Event
Mon 3/11	Assembly Wrap
Tue/Wed	Review Lab
Wed 3/13	Review
Thu 3/14	A3 Due
Fri 3/15	Exam 2
3/18-3/24	Spring Break

Exam 2 Review Wed

Expect a practice exam

Assignment 3: Questions?

Exercise: All Models are Wrong...

- ▶ Rule #1: The Doctor Lies
- ▶ Below is our original model for memory layout of C programs
- ▶ Describe what is **wrong** based on x86-assembly
- ▶ Will all variables have a position in the stack?
- ▶ What else is on the stack / control flow info?
- ▶ What registers are likely used?

```
9: int main(...){
10:   int x = 19;
11:   int y = 31;
+<-12: swap(&x, &y);
| 13:   printf("%d %d\n",x,y);
| 14:   return 0;
V 15: }
```

STACK: Caller main(), prior to swap()

FRAME	ADDR	NAME	VALUE
-----+-----+-----+-----			
main()	#2048	x	19
line:12	#2044	y	31
-----+-----+-----+-----			

```
|
| 18: void swap(int *a,int *b){
+>-19:   int tmp = *a;
20:   *a = *b;
21:   *b = tmp;
22:   return;
23: }
```

STACK: Callee swap() takes control

FRAME	ADDR	NAME	VALUE	
-----+-----+-----+-----				
main()	#2048	x	19	<-+
line:12	#2044	y	31	<- +
-----+-----+-----+-----				
swap()	#2040	a	#2048	---+
line:19	#2036	b	#2044	----+
	#2032	tmp	?	

Answers: All Models are Wrong, Some are Useful

```

    9: int main(...){
    10:   int x = 19;
    11:   int y = 31;
+-->12:   swap(&x, &y);
|  13:   printf("%d %d\n",x,y);
|  14:   return 0;
V  15: }
|
|  18: void swap(int *a,int *b){
+-->19:   int tmp = *a;
    20:   *a = *b;
    21:   *b = tmp;
    22:   return;
    23: }
```

STACK: Callee swap() takes control

FRAME	ADDR	NAME	VALUE
-----+-----+-----+-----			
main()	#2048	x	19
	#2044	y	31
-----+-----+-----+-----			
swap()	#2036	rip	L12
-----+-----+-----+-----			

REGS as swap() starts

REG	VALUE	NOTE
-----+-----+-----		
rdi	#2048	for *a
rsi	#2044	for *b
rax	?	for tmp
rip	L19	line in swap

- ▶ main() must have stack space for locals passed by address
- ▶ swap() needs no stack space for arguments: in registers
- ▶ Return address is old value of rip register
- ▶ Mostly don't need to think at this level of detail but **can be useful in some situations**

Data In Assembly

Arrays

Usually: $\text{base} + \text{index} \times \text{size}$

```
arr[i] = 12;  
movl $12, (%rdi,%rsi,4)
```

```
int x = arr[j];  
movl (%rdi,%rcx,4), %r8
```

- ▶ Array starting address often held in a register
- ▶ Index often in a register
- ▶ Compiler inserts appropriate size (1,2,4,8)

Structs

Usually $\text{base} + \text{offset}$

```
typedef struct {  
    int i; short s;  
    char c[2];  
} foo_t;  
foo_t *f = ...;
```

```
short sh = f->s;  
movw 4(%rdi), %si
```

```
f->c[i] = 'X';  
movb $88, 6(%rdi,%rax)
```

General Cautions on Structs

Struct Layout: Compiler calculates

- ▶ Ordering of fields struct in memory
- ▶ Padding between/after fields for alignment
- ▶ Total struct size

Struct Layout Algorithms

- ▶ Baked into compiler
- ▶ **May change from compiler to compiler**
- ▶ May change through history of compiler

Structs in Mem/Regs

- ▶ Stack structs spread across several registers
- ▶ Don't need a struct on the stack at all in some cases (just like don't need local variables on stack)
- ▶ Struct arguments packed into 1+ registers

Stay Insulated

- ▶ Programming in C insulates you from all of this
- ▶ Feel the **warmth** of gcc's abstraction blanket

Security Risks in C

Buffer Overflow Attacks

- ▶ No default bounds checking in C: Performance favored over safety
- ▶ Allows classic security flaws:

```
char buf[1024];
printf("Enter you name:");
fscanf(file,"%s",buf); // BAD
// or
gets(buf);              // BAD
// my name is 1500 chars
// long, what happens?
```
- ▶ For data larger than buffer, begin overwriting other parts of the stack
 - ▶ Clobber return addresses
 - ▶ Insert executable code and run it

Counter-measures

- ▶ **Stack protection** is default in gcc in the modern era
- ▶ Inserts "canary" values on the stack near return address
- ▶ Prior to function return, checks that canaries are unchanged
- ▶ **Stack (start) randomized** by kernel making address of functions difficult to predict ahead of time
- ▶ Kernel may also vary virtual memory address as well
- ▶ Disable protections at your own risk

Sample Buffer Overflow Code

```
#include <stdio.h>
void never(){
    printf("This should never happen\b");
    return;
}
int main(){
    union {long addr; char str[9];} never_info;
    never_info.addr = (long) never;
    never_info.str[8] = '\\0';

    printf("Address of never: %0p\\n",never_info.addr);
    printf("Address as string: %s\\n",never_info.str);

    printf("Enter a string: ");
    char buf[4];
    fscanf(stdin,"%s",buf);
    // By entering the correct length of string followed by the ASCII
    // representation of the address of never(), one might be able to
    // get that function to run (on windows...)

    printf("You entered: %s\\n",buf);
    return 0;
}
```


Floating Point Operations

- ▶ The original Intel Chips 8086 **didn't have floating point ops**
- ▶ Had to buy a co-processor, Intel 8087, to add FP ops
- ▶ Modern CPUs ALL have FP ops but they feel separate from the integer ops: FP Unit versus AL Unit

FP "Media" Registers

256-bits	128-bits	Use
%ymm0	%xmm0	FP Arg 1/ Ret
%ymm1	%xmm2	FP Arg 2
...
%ymm7	%xmm7	FP Arg 8
%ymm8	%xmm8	Caller Save
...
%ymm15	%xmm15	Caller Save

- ▶ Can be used as "scalars" - single values but...
- ▶ `xmmI` is 128 bits big holding
 - ▶ 2 64-bit FP values OR
 - ▶ 4 32-bit FP values
- ▶ `ymmI` doubles this

Instructions

- ▶ Usually 3 operands:
 $C = B \text{ op } A$
- ▶ Ex: Subtraction `vsubsd`, with `d` for 64-bit double
`# xmm0 = xmm2 - xmm4`
`vsubsd %xmm2,%xmm4,%xmm0`
- ▶ 3-operands common in modern assembly
- ▶ Can operate on single values or "vectors" of packed values

Floating Point and ALU Conversions

- ▶ Recall that bit layout of Integers and Floating Point numbers are quite different (**how?**)
- ▶ Leads to a series of assembly instructions to interconvert between types

```
# int eax = ...;  
# double xmm0 = (double) eax;  
vcvtsi2sd      %eax,%xmm0,%xmm0
```

```
# double xmm1 = ...  
# long rcx = (int) xmm1;  
vcvttsd2siq    %xmm1,%rcx
```

- ▶ These are non-trivial conversions: 5-cycle latency (delay) before completion, can have a performance impact on code which does conversions