

# CSCI 2021: Program Performance Micro-Optimizations

Chris Kauffman

*Last Updated:  
Tue Apr 16 09:25:06 CDT 2019*

# Logistics

Date	Event
Fri 4/12	Lec: Micro-opts
Mon 4/15	Lec: Micro-opts,
Tue 4/16	<b>A4 DUE</b>
Wed 4/17	Lec/Lab11: Review
Fri 4/19	Exam 3

## Reading Bryant/O'Hallaron

Ch 5: Optimizing Program  
Performance

## Goals

- ▶ What to Optimize First
- ▶ First-best Optimizations
- ▶ Micro-optimization  
Techniques

## Assignment 4

- ▶ Questions?

## Lab 10


- ▶ Comparing  
micro-optimizations
- ▶ Memory optimizations

# Caution: Should I Optimize?

- ▶ Optimizing program execution time usually costs human time
- ▶ Human time is valuable, don't waste it
- ▶ Determine if there is a **NEED** to optimize
- ▶ **Benchmark** your code - if it is fast enough, move on
- ▶ If not fast enough, use a **profiler** to determine where your efforts are best spent
- ▶ **Never sacrifice correctness** for speed

First make it **work**,  
then make it **right**,  
then make it **fast**.

- Kent Beck



# What to Optimize First

In order of impact

1. Algorithms and Data Structure Selection
2. Elimination of unneeded work/hidden costs
3. Memory Utilization
4. **Micro-optimizations**

**“Premature optimization is the root of all evil” - Donald Knuth**



Programmers waste enormous amounts of time thinking about, or worrying about, the speed of noncritical parts of their programs, and these attempts at efficiency actually have a strong negative impact when debugging and maintenance are considered. We should forget about small efficiencies, say about 97% of the time: *premature optimization is the root of all evil.* **Yet we should not pass up our opportunities in that critical 3%.**

– Donald Knuth

## Exercise: Optimize This

- ▶ Prema Turopt is tasked by her boss to optimize the performance function `get_min()`
- ▶ The current version of the function code looks like the code to the right.
- ▶ Prema immediately jumps to the code for `bubble_sort()` and alters the code to enable better processor pipelining.
- ▶ This leads to a 3% improvement in speed.

```
1  int get_min(storage_t *st){
2      int *arr =
3          malloc(sizeof(int)*get_size(st));
4
5      for(int i=0; i<get_size(st); i++){
6          arr[i] = get_element(st,i);
7      }
8
9      bubble_sort(arr, get_size(st));
10
11     int ans = arr[0];
12     free(arr);
13     return ans;
14 }
15
```

**Suggest several alternatives** that Prema should have explored

# Answers: Optimize This

1. Don't use bubblesort:  $O(N^2)$ .  
Use an  $O(N \log N)$  sort like Quicksort, Heapsort, Mergesort
2. Why sort at all? Determine the minimum element with the "get" loop.
3. What is the cost of `get_element()` and `get_size()`? Is there a more efficient iterator or array-extraction mechanism?
4. What data structure is used in `storage_t`? If it is already sorted such as a binary search tree or binary heap, there may be a more efficient way to determine the minimum element.

```
1  int get_min(storage_t *st){
2      int *arr =
3          malloc(sizeof(int)*get_size(st));
4
5      for(int i=0; i<get_size(st); i++){
6          arr[i] = get_element(st,i);
7      }
8
9      bubble_sort(arr, get_size(st));
10
11     int ans = arr[0];
12     free(arr);
13     return ans;
14 }
```

5. Might be able to arrange for minimum elements to be tracked automatically in the data structure for `storage_t` if other code can be modified -  $O(1)$  `get_min()` method.

## Exercise: Eliminating Unnecessary Work

```
void lower1(char *s) {  
    for (long i=0; i < strlen(s); i++){  
        if (s[i] >= 'A' && s[i] <= 'Z'){  
            s[i] -= ('A' - 'a');  
        }  
    }  
}
```

```
void lower2(char *s) {  
    long len = strlen(s);  
    for (long i=0; i < len; i++){  
        if (s[i] >= 'A' && s[i] <= 'Z'){  
            s[i] -= ('A' - 'a');  
        }  
    }  
}
```

- ▶ Bryant/O'Hallaron Figure 5.7
- ▶ Two versions of a lower-casing function
- ▶ Lowercase by subtracting off constant for uppercase characters: alters ASCII code
- ▶ Examine them to determine differences
- ▶ Project speed differences and **why one will be faster**


# Answers: Eliminating Unnecessary Work

- ▶ `strlen()` is  $O(N)$ : searches for `\0` character in `for()` loop
- ▶ Don't loop with it if possible

```
void lower1(char *s) {  
    for (long i=0; i < strlen(s); i++){  
        if (s[i] >= 'A' && s[i] <= 'Z'){  
            s[i] -= ('A' - 'a');  
        }  
    }  
}
```

```
long strlen(char *s) {  
    long len = 0;  
    while(s[len] != '\0'){  
        len++;  
    }  
    return len;  
}
```

```
void lower2(char *s) {  
    long len = strlen(s);  
    for (long i=0; i < len; i++){  
        if (s[i] >= 'A' && s[i] <= 'Z'){  
            s[i] -= ('A' - 'a');  
        }  
    }  
}
```





## Exercise: Allocation and Hidden Costs

Consider the following **Java** code

```
public class StringUtils{
    public static
    String repString(String str, int reps)
    {
        String result = "";
        for(int i=0; i<reps; i++){
            result = result + str;
        }
        return result;
    }
}
```

- ▶ Give a Big-O estimate for the runtime
- ▶ Give a Big-O estimate for the memory overhead

# Answers: Allocation and Hidden Costs

- ▶ Strings are **immutable** in Java (Python, many others)
- ▶ Each iteration must
  - ▶ **allocate** new memory for a new string sized `result.length + str.length`
  - ▶ Copy result to the first part
  - ▶ Copy str to the second part
- ▶ Leads to  $O(N^2)$  complexity
- ▶ Much worse memory usage: as much as  $O(N^2)$  wasted memory for garbage collector to clean up

```
public class StringUtils{
    public static
    String repString(String str, int reps)
    {
        String result = "";
        for(int i=0; i<reps; i++){
            result = result + str;
        }
        return result;
    }

    // Efficient version
    public static
    String repString2(String str, int reps)
    {
        StringBuilder result =
            new StringBuilder();
        for(int i=0; i<reps; i++){
            result.append(str);
        }
        return result.toString();
    }
}
```

## Exercise: Do Memory References Matter?

```
void sum_range1(int start,
               int stop,
               int *ans)
{
    *ans = 0;
    for(int i=start; i<stop; i++){
        *ans += i;
    }
}
```

```
void sum_range2(int start,
               int stop,
               int *ans)
{
    int sum = 0;
    for(int i=start; i<stop; i++){
        sum += i;
    }
    *ans = sum;
}
```

- ▶ What is the primary difference between the two routines above?
- ▶ What effect if any will this have on runtime?

# Answers: Do Memory References Matter?

```
void sum_range1(int start,
                int stop,
                int *ans)
{
    *ans = 0;
    for(int i=start; i<stop; i++){
        *ans += i;
    }
}
```

sum\_range1() makes repeated memory references

```
void sum_range2(int start,
                int stop,
                int *ans)
{
    int sum = 0;
    for(int i=start; i<stop; i++){
        sum += i;
    }
    *ans = sum;
}
```

sum\_range2() uses a local variable with only a couple memory references

- ▶ Must determine if memory references matter for performance
- ▶ Guesses?

# Memory References Matter, Compiler May Change Them

```
lila> gcc -Og sum_range.c # No opt
lila> ./a.out 0 1000000000
sum_range1: 1.9126e+00 secs
sum_range2: 2.6942e-01 secs
```

- ▶ Minimal optimizations
- ▶ Memory reference definitely matters

```
lila> gcc -O1 sum_range.c # Opt plz
lila> ./a.out 0 1000000000
sum_range1: 2.8972e-01 secs
sum_range2: 2.7569e-01 secs
```

- ▶ Observe code differences between -Og and -O1
- ▶ Why is performance improved so much?

```
### Compiled with -Og: minimal opt
sum_range1:
    movl    $0, (%rdx)    # write to memory
    jmp     .LOOPTHOP
.BODY:
    addl    %edi, (%rdx)  # memory write
    addl    $1, %edi      # in loop
.LOOPTHOP:
    cmpl    %esi, %edi
    jl      .BODY
    ret
```

```
### Compiled with -O1: some opt
sum_range1:
    movl    $0, (%rdx)    # mem write
    cmpl    %esi, %edi
    jge     .END
    movl    $0, %eax      # 0 to reg
.LOOP:
    addl    %edi, %eax     # add to reg
    addl    $1, %edi      # no mem ref
    cmpl    %edi, %esi
    jne     .LOOP
    movl    %eax, (%rdx)  # write at end
.END:
    ret
```

# Dash-O! Compiler Optimizes for You

- ▶ gcc like, many compilers, can perform some memory and micro-optimizations for you
- ▶ Series of -O<sub>X</sub> options to enable different techniques
- ▶ We will use -O<sub>g</sub> at times to disable many optimizations
  - ▶ -O<sub>g</sub>: Optimize debugging. ... optimization level of choice for the standard edit-compile- debug cycle
- ▶ Individual optimizations can be enabled and disabled
- ▶ -O or -O<sub>1</sub>: Optimize. Optimizing compilation takes somewhat more time, and a lot more memory for a large function.  
With -O, the compiler tries to reduce code size and execution time, without performing any optimizations that take a great deal of compilation time.
- ▶ -O<sub>2</sub>: Optimize even more. GCC performs nearly all supported optimizations that do not involve a space-speed tradeoff. As compared to -O, this option increases both compilation time and the performance of the generated code.
- ▶ -O<sub>3</sub>: Optimize yet more. -O<sub>3</sub> turns on all optimizations specified by -O<sub>2</sub> and also...
- ▶ -Ofast: Disregard strict standards compliance. (!)

# Compiler Optimizations

-O turns on the following optimization flags:

```
-fauto-inc-dec -fbranch-count-reg -fcombine-stack-adjustments  
--fcompare-elim fccprop-registers -fdce -fdefer-pop -fdelayed-branch  
--fdse -fforward-propagate fguess-branch-probability -fif-conversion2  
--fif-conversion finline-functions-called-once -fipa-pure-const  
--fipa-profile -fipa-reference fmerge-constants -fmove-loop-invariants  
--freorder-blocks -fshrink-wrap fshrink-wrap-separate  
--fsplit-wide-types -fssa-backprop -fssa-phiopt -ftree-bit-ccp  
-ftree-ccp -ftree-ch -ftree-coalesce-vars -ftree-copy-prop -ftree-dce  
-ftree-dominator-opts -ftree-dse -ftree-forwprop -ftree-fre  
--ftree-phi-prop -ftree-sink -ftree-slsr -ftree-sra -ftree-pta  
--ftree-ter -funit-at-a-time
```

- ▶ Some combination of these enables `sum_range2()` to fly as fast as `sum_range1()`
- ▶ We will look at some “by-hand” versions of these optimizations but whenever possible, let the compiler do it

## Exercise: Loop Unrolling

- ▶ Have seen copying loop iterations manually *may* lead to speed gains
- ▶ **Why?** Which of the following unrolled versions of `sum_rangeX()` seems fastest?
- ▶ Why the **second loop** in `sum_rangeB()` and `sum_rangeC()`?

```
1 void sum_rangeA(long stop, long *ans){
2     long sum=0, i;
3     for(i=0; i<stop; i++){
4         sum += i+0;
5     }
6     *ans = sum;
7 }
8
```

```
9 void sum_rangeB(long stop, long *ans){
10     long sum = 0, i;
11     for(i=0; i<stop-3; i+=3){
12         sum += (i+0);
13         sum += (i+1);
14         sum += (i+2);
15     }
16     for(; i<stop; i++){
17         sum += i;
18     }
19     *ans = sum;
20 }
21
22 void sum_rangeC(long stop, long *ans){
23     long sum0=0, sum1=0, sum2=0, i;
24     for(i=0; i<stop-3; i+=3){
25         sum0 += (i+0);
26         sum1 += (i+1);
27         sum2 += (i+2);
28     }
29     for(; i<stop; i++){
30         sum0 += i;
31     }
32     *ans = sum0 + sum1 + sum2;
33 }
```



## Exercise: Loop Unrolling

### Expectations

Version	Notes	Performance
sum_rangeA()	Not unrolled	Baseline
sum_rangeB()	Unroll x3, same destinations for sum	Less good
sum_rangeC()	Unroll x3, different destinations sum add	Expected Best

### Actual Performance

```
apollo> gcc -Og unroll.c
apollo> ./a.out 1000000000
sum_rangeA: 1.0698e+00 secs
sum_rangeB: 6.2750e-01 secs
sum_rangeC: 6.2746e-01 secs

phaedrus> ./a.out 1000000000
sum_rangeA: 2.8913e-01 secs
sum_rangeB: 5.3285e-01 secs
sum_rangeC: 2.6774e-01 secs
```

### Unrolling is Unpredictable

- ▶ Performance Gains vary from one compiler+processor to another
- ▶ All unrolling requires **cleanup loops** like those in the B/C versions: add on remaining elements

# GCC Options to Unroll

- ▶ gcc has options to unroll loops during optimization
- ▶ Unrolling has unpredictable performance implications so unrolling is **not enabled** for -O1, -O2, -O3
- ▶ Can manually enable it with compiler options like -funroll-loops to check for performance bumps

```
apollo> gcc -Og unroll.c
apollo> ./a.out 1000000000
sum_rangeA: 1.0698e+00 secs
sum_rangeB: 6.2750e-01 secs
sum_rangeC: 6.2746e-01 secs
```

```
apollo> gcc -Og -funroll-loops unroll.c
apollo> ./a.out 1000000000
sum_rangeA: 7.0386e-01 secs    # loop unrolled by compiler
sum_rangeB: 6.2802e-01 secs
sum_rangeC: 6.2797e-01 secs
```

```
apollo> gcc -Og -funroll-loops -fvariable-expansion-in-unroller unroll.c
apollo> ./a.out 1000000000
sum_rangeA: 5.2711e-01 secs    # unroll + multiple intermediates used
sum_rangeB: 6.2759e-01 secs
sum_rangeC: 6.2750e-01 secs
```

```
apollo> gcc -O3 unroll.c
apollo> ./a.out 1000000000
sum_rangeA: 9.4124e-01 secs
sum_rangeB: 4.1833e-01 secs
sum_rangeC: 4.1832e-01 secs
# manual unroll + compiler opt
```

# Do Conditionals Matter?

Consider two examples of adding even numbers in a range

```
1  // CONDITION version
2  long sum_evensA(long start, long stop){
3      long sum=0;
4      for(int i=start; i<stop; i++){
5          if((i & 0x01) == 0){
6              sum += i;
7          }
8      }
9      return sum;
10 }
11 // STRAIGHT-LINE version
12 long sum_evensB(long start, long stop){
13     long sum=0;
14     for(int i=start; i<stop; i++){
15         int odd = i & 0x01;
16         int even_mask = odd - 1;
17         // 0x00000000 for odd
18         // 0xFFFFFFFF for even
19         sum += even_mask & i;
20     }
21     return sum;
22 }
```

Timings for these two are shown below at two levels of optimization.

```
lila> gcc -Og condloop.c
lila> a.out 0 400000000
sum_evensA: 1.1969e+00 secs
sum_evensB: 2.8953e-01 secs
# 4x speedup
```

```
lila> gcc -O3 condloop.c
lila> a.out 0 400000000
sum_evensA: 2.3662e-01 secs
sum_evensB: 9.6242e-02 secs
# 2x speedup
```

Message is simple: **eliminate conditionals** whenever possible to improve performance

## Exercise: Row Sums with Function v Macro

What is the difference between these two row sum functions?

```
1  int mget(matrix_t mat,
2      int i, int j)
3  {
4      return
5      mat.data[i*mat.cols + j];
6  }
7  int vset(vector_t vec,
8      int i, int x)
9  {
10     return vec.data[i] = x;
11 }
12 void row_sumsA(matrix_t mat,
13     vector_t sums)
14 {
15     for(int i=0; i<mat.rows; i++){
16         int sum = 0;
17         for(int j=0; j<mat.cols; j++){
18             sum += mget(mat,i,j);
19         }
20         vset(sums, i, sum);
21     }
22 }
```

```
1  #define MGET(mat,i,j) \
2      ((mat).data[((i)*((mat).cols)) + (j)])
3
4
5
6
7  #define VSET(vec,i,x) \
8      ((vec).data[(i)] = (x))
9
10
11
12 void row_sumsB(matrix_t mat,
13     vector_t sums)
14 {
15     for(int i=0; i<mat.rows; i++){
16         int sum = 0;
17         for(int j=0; j<mat.cols; j++){
18             sum += MGET(mat,i,j);
19         }
20         VSET(sums, i, sum);
21     }
22 }
```

## Answers: Row Sums with Function v Macro

- ▶ `row_sumsA()` uses standard function calls to retrieve elements
- ▶ `row_sumsB()` uses **macros** to do the element retrieval
- ▶ A macro is a textual expansion done by the **preprocessor**: insert the literal text associated with the macro
- ▶ See macro results with `gcc -E rowsums.c` which stops after preprocessor step (early)
- ▶ Function calls cost some operations but not many
- ▶ Function calls **prevent optimization across boundaries**
- ▶ Cannot pipeline effectively when jumping around, using registers for arguments, restoring registers, etc
- ▶ Macros can alleviate this but they are a **pain** to write and notoriously buggy
- ▶ Better to let the compiler do this for us

# Inlining Functions/Procedures

- ▶ **Function Inlining** inserts the body of a function where it would have been called
- ▶ Turned on fully partially at -O2 and fully at -O3
- ▶ Enables other optimizations blocked by function boundaries
- ▶ Can only be done if source code (C file) for function is available
- ▶ Like loop unrolling, function inlining has trade-offs
  - ▶ Enables pipelining
  - ▶ More predictable control
  - ▶ More register pressure
  - ▶ Increased code size

```
> FILES="rowsums.c matvec_util.c"
> gcc -Og $FILES
> ./a.out 8000 8000
row_sumsA: 1.3109e-01 secs
row_sumsB: 4.0536e-02 secs
```

```
> gcc -Og -finline-small-functions $FILES
> ./a.out 8000 8000
row_sumsA: 7.4349e-02 secs
row_sumsB: 4.2682e-02 secs
```

```
> gcc -O3 $FILES
> ./a.out 8000 8000
row_sumsA: 2.1974e-02 secs
row_sumsB: 2.0820e-02 secs
```

- ▶ Inlining typically most effective for small functions (getters/setters)

# Profilers: gprof and Friends

- ▶ **Profiler:** a tool that monitors code execution to enable performance optimizations
- ▶ gprof is stock on Linux systems, interfaces with gcc
- ▶ Compile with profiling options: `gcc -pg`
- ▶ Run code to produce data file
- ▶ Examine with gprof
- ▶ **Note:** gcc version 6 and 7 contain a bug requiring use of `-no-pie` option, not a problem on apollo

```
# Compile
# -pg : instrument code for profiling
# -no-pie : bug fix for new-ish gcc's
> gcc -pg -no-pie -g -Og -o unroll unroll.c

> ls
unroll  unroll.c

> ./unroll 1000000000
sum_rangeA: 2.9401e-01 secs
sum_rangeB: 5.3164e-01 secs
sum_rangeC: 2.6574e-01 secs

# gmon.out now created with timing info
> ls
gmon.out  unroll  unroll.c

> file gmon.out
gmon.out: GNU prof performance data

> gprof -b unroll
... output on next slide ...
```

## gprof output for unroll

```
> gprof -b unroll
```

```
Flat profile:
```

```
Each sample counts as 0.01 seconds.
```

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
50.38	0.54	0.54	1	544.06	544.06	sum_rangeB
26.12	0.83	0.28	1	282.11	282.11	sum_rangeA
24.26	1.09	0.26	1	261.95	261.95	sum_rangeC

### Call graph

index	% time	self	children	called	name
[1]	100.0	0.00	1.09		main [1]
		0.54	0.00	1/1	sum_rangeB [2]
		0.28	0.00	1/1	sum_rangeA [3]
		0.26	0.00	1/1	sum_rangeC [4]
-----					
[2]	50.0	0.54	0.00	1/1	main [1]
		0.54	0.00	1	sum_rangeB [2]
-----					
[3]	25.9	0.28	0.00	1/1	main [1]
		0.28	0.00	1	sum_rangeA [3]
-----					
[4]	24.1	0.26	0.00	1/1	main [1]
		0.26	0.00	1	sum_rangeC [4]
-----					



## gprof Example: Dictionary Application

```
> ./dictionary < craft-67.txt
... Total time = 0.829561 seconds
> gprof -b dictionary
```

	%	cumulative	self		self	total	
time		seconds	seconds	calls	ms/call	ms/call	name
50.07	0.18	0.18	1	180.25	180.25	sort_words	
19.47	0.25	0.07	463016	0.00	0.00	find_ele_rec	
13.91	0.30	0.05	2862749	0.00	0.00	Strlen	
8.34	0.33	0.03	463016	0.00	0.00	lower1	
2.78	0.34	0.01	463017	0.00	0.00	get_token	
2.78	0.35	0.01	463016	0.00	0.00	h_mod	
2.78	0.36	0.01	20451	0.00	0.00	save_string	
0.00	0.36	0.00	463017	0.00	0.00	get_word	
0.00	0.36	0.00	463016	0.00	0.00	insert_string	
0.00	0.36	0.00	20451	0.00	0.00	new_ele	
0.00	0.36	0.00	7	0.00	0.00	add_int_option	
0.00	0.36	0.00	1	0.00	0.00	add_string_option	
0.00	0.36	0.00	1	0.00	0.00	init_token	
0.00	0.36	0.00	1	0.00	0.00	new_table	
0.00	0.36	0.00	1	0.00	0.00	parse_options	
0.00	0.36	0.00	1	0.00	0.00	show_options	
0.00	0.36	0.00	1	0.00	360.50	word_freq	

## gprof Example Cont'd: Dictionary Application

```
> ./dictionary < craft-67.txt      ## After upgrading sort_words() to qsort()
... Total time = 0.624172 seconds
> gprof -b dictionary
```

time	% cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
60.08	0.12	0.12	463016	0.00	0.00	find_ele_rec
15.02	0.15	0.03	2862749	0.00	0.00	Strlen
10.01	0.17	0.02	463016	0.00	0.00	lower1
5.01	0.18	0.01	463017	0.00	0.00	get_token
5.01	0.19	0.01	463016	0.00	0.00	h_mod
5.01	0.20	0.01	20451	0.00	0.00	save_string
0.00	0.20	0.00	463017	0.00	0.00	get_word
0.00	0.20	0.00	463016	0.00	0.00	insert_string
0.00	0.20	0.00	20451	0.00	0.00	new_ele
0.00	0.20	0.00	8	0.00	0.00	match_length
0.00	0.20	0.00	7	0.00	0.00	add_int_option
0.00	0.20	0.00	1	0.00	0.00	add_string_option
0.00	0.20	0.00	1	0.00	0.00	find_option
0.00	0.20	0.00	1	0.00	0.00	init_token
0.00	0.20	0.00	1	0.00	0.00	new_table
0.00	0.20	0.00	1	0.00	0.00	parse_options
0.00	0.20	0.00	1	0.00	0.00	show_options
0.00	0.20	0.00	1	0.00	0.00	sort_words ** was 0.18 **
0.00	0.20	0.00	1	0.00	200.28	word_freq